

Technical Framework for Subtree Selection and Node Weighting in Automated Prerequisite Course Planning

1 Introduction

For a given target course, prerequisites can be represented as a complex tree of requirements. The primary objective of this work is to define a technical framework for navigating this structure to find an optimal path for the student.

Specifically, given a prerequisite tree, the goal is to select an optimal directed subtree rooted at the target course. This subtree must satisfy all logical requirements (i.e., AND/OR conditions). The selection process is guided by a weighting system that reflects course quality and user preferences, aiming to maximize the total weight of the selected courses while promoting course reuse to ensure an efficient academic path.

2 Methodology

2.1 Prerequisite Graph Model

We model the prerequisite structure as a rooted Directed Acyclic Graph (DAG), denoted $G = (V, E)$, with the following properties:

- The vertex set V consists of two types of nodes:
 - **Course nodes**, which represent atomic courses, possibly with a minimum grade requirement.
 - **Logic nodes**, which represent either an AND (\wedge) or OR (\vee) condition over their children.
- The root of the DAG is the target course for which the prerequisite path is being planned.

- Edges in E represent dependencies, connecting parent nodes to their required sub-requirements (children).

2.2 Optimal Subtree Selection

The core of the methodology is a recursive function, $\text{SelectSubtree}(T, S)$, designed to traverse the graph and identify the optimal set of courses. Here, T is the current node (either course or logic) under consideration, and S is the set of courses already selected, which is used to handle dependencies and prevent duplicates. The function is detailed in Algorithm 1.

Algorithm 1 Optimal Subtree Selection for Prerequisite Satisfaction

```

1: function SELECTSUBTREE( $T, S$ )
2:   if  $T$  is a course node then
3:     if  $T \notin S$  then
4:        $P \leftarrow \{T\}$ 
5:        $S \leftarrow S \cup \{T\}$ 
6:       for all prerequisite  $Q$  of  $T$  do
7:          $P \leftarrow P \cup \text{SELECTSUBTREE}(Q, S)$ 
8:       end for
9:       return  $P$ 
10:    else
11:      return  $\emptyset$ 
12:    end if
13:  else if  $T$  is an AND-node then
14:     $P \leftarrow \emptyset$ 
15:    for all child  $C$  of  $T$  do
16:       $P \leftarrow P \cup \text{SELECTSUBTREE}(C, S)$ 
17:    end for
18:    return  $P$ 
19:  else if  $T$  is an OR-node then
20:     $P^* \leftarrow \text{argmin}_P \text{Cost}(P)$  over all  $P = \text{SELECTSUBTREE}(C, S)$  for
    children  $C$  of  $T$ 
21:    return  $P^*$ 
22:  end if
23: end function

```

The cost function, $\text{Cost}(P)$, used at OR-nodes, is initially defined as the negative sum of weights of the courses in a potential path P : $\sum_{v \in P} -y_v$. This serves to maximize the total path weight. This function is expanded

in Section 4.

2.3 Algorithmic Properties

- **Feasibility:** The algorithm guarantees a feasible path by only omitting courses that are already in the selected set S , ensuring that all dependencies are met.
- **Optimality at OR-nodes:** At each OR-node, the branch that yields the optimal cost (e.g., maximum weight, minimum courses) is chosen based on the defined cost function.
- **Completeness at AND-nodes:** At each AND-node, all branches are traversed, and their resulting course sets are unified to ensure all requirements are satisfied.
- **Theoretical Basis:** This approach frames course selection as a dynamic programming problem on a tree, which can be seen as a generalization of the AND/OR pathfinding problem in graphs.

3 Node Weighting Algorithm

Each course node is assigned a weight y_i to quantify its desirability based on historical student feedback and user-specified preferences.

3.1 Parameters

For each course i , we define the following parameters:

- x_1 : liked score $[0, 100]$
- x_2 : easy score $[0, 100]$
- x_3 : useful score $[0, 100]$
- r_i : number of student ratings

User preferences are captured by a weight vector $\boldsymbol{\beta} = (\beta_1, \beta_2, \beta_3)$, where $\beta_k \geq 0$ and $\sum_k \beta_k = 1$. Predefined profiles for $\boldsymbol{\beta}$ include:

- **Focus on likeness:** $(0.7, 0.15, 0.15)$
- **Focus on easiness:** $(0.15, 0.7, 0.15)$
- **Focus on usefulness:** $(0.15, 0.15, 0.7)$
- **Balanced:** $(0.4, 0.3, 0.3)$

3.2 Reliability Scaling Factor

To account for the confidence in course ratings, a reliability scaling factor, λ_i , is introduced, which is monotonic in the number of ratings r_i . It can be defined using discrete thresholds:

$$\lambda_i = \begin{cases} 1.10 & \text{if } r_i > 100 \\ 1.00 & \text{if } 50 \leq r_i \leq 100 \\ 0.9 & \text{if } r_i < 50 \end{cases}$$

Alternatively, a continuous scaling function can be used for finer-grained adjustments:

$$\lambda_i = 1 + \alpha \cdot \left(\frac{r_i - \bar{r}}{\sigma_r} \right), \quad \text{clamped to } [0.95, 1.05]$$

where $\bar{r} = 32.8743$ and $\sigma_r = 67.55$ are the mean and standard deviation of ratings across all courses, and $\alpha \approx 0.15$ is a small constant.

3.3 Null Value Handling

Missing rating values (x_1, x_2, x_3) are handled as follows:

- If **one or two** scores are null, they are imputed using the column median (or mean) over all courses. Alternatively, the remaining β_k weights can be re-normalized to sum to 1.
- If **all three** scores are null, they are assigned a low value to penalize the course while keeping it as a viable option:

$$x_k^* = \gamma \cdot \min_{\text{all non-null values}} x_{j,\ell}$$

for $k = 1, 2, 3$, with a penalty factor $\gamma \in [0.3, 0.7]$ (e.g., $\gamma = 0.5$). For such courses, λ_i is set to a penalty value like 0.97 to further reflect their unreliability.

3.4 Node Weight Calculation

The final weight y_i for a course is calculated as:

$$s_i = \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3$$

$$y_i = \min(\lambda_i s_i, 100)$$

The min operation ensures the weight does not exceed the maximum possible score of 100.

3.5 Edge Cases

- Courses with all rating scores imputed will have a small, non-zero weight. This ensures they are only selected if no other paths are available.
- If all available paths require traversing a node with imputed scores, the algorithm will still return the path with the highest possible total weight.

4 Global Reuse Preference (Depth-Aware New-Course Penalty)

The initial objective function encourages high-weight courses but does not explicitly favor solutions that reuse courses across different prerequisite branches. To promote efficient paths, we introduce a *new-course penalty*: a depth-aware cost added only when a course *ID* appears in the path for the first time.

4.1 Penalty Definition

Let S be the set of selected course IDs so far. For a course node v at depth d (root at $d = 0$), the base cost is

$$\text{baseCost}(v) = 1 - \frac{y_v}{100}.$$

A depth-aware penalty is charged only on first inclusion:

$$\text{newCoursePenalty}(v, S, d) = \begin{cases} \lambda(d) & v \notin S, \\ 0 & v \in S, \end{cases} \quad \lambda(d) = \frac{\lambda_0}{1+d}, \quad \lambda_0 \in [0.1, 0.4] \text{ (default 0.2)}.$$

The node cost is then

$$\text{cost}(v, S, d) = \text{baseCost}(v) + \text{newCoursePenalty}(v, S, d).$$

This is equivalent to a “reuse award” formulation up to a constant shift; we use the new-course penalty because it is simpler to implement and reason about.

4.2 Selection with Global Reuse

The recursive selection passes (S, d) so that OR-branches are evaluated from the same state:

- **Course node:** Return $(\text{cost}(v, S, d), S \cup \{v\})$.
- **AND node:** Sum child costs and accumulate S through the children.
- **OR node:** For each child C , evaluate with the same incoming S and pick the minimal-cost branch (stable tie-break by order).

Base-cost rule on reuse. When a course reappears in another branch, the base cost $1 - y_v/100$ and the penalty are charged *only* on the first inclusion of its ID in S . Later reuses add zero cost. This ensures that reusing a course strictly reduces $\mathcal{C}(P)$ relative to introducing a new course, all else equal.

Notes.

- The node weight y_v continues to guide the selection towards user-preferred courses, while the reuse penalty refines the selection to favor more efficient paths.
- The penalty schedule $\lambda_u(d)$ is a tunable hyperparameter and can be replaced by any monotonically decreasing function of depth.

4.3 Objective Function

The overall objective is to find a valid prerequisite path P that minimizes the total cost function $\mathcal{C}(P)$, which combines the base cost of courses with penalties for introducing new ones.

$$\min_P \mathcal{C}(P) = \min_P \sum_{v \in P} \text{cost}(v, S_v, d_v)$$

where for each node $v \in P$:

- d_v is its depth, and S_v is the set of courses selected prior to visiting v .
- $\text{cost}(v, S_v, d_v) = (1 - y_v/100) + \text{uniquePenalty}(v, S_v, d_v)$.

The **SelectSubtree** function acts as a recursive procedure to find the path P that minimizes this sum by making locally optimal decisions at OR-nodes while propagating the global set of selected courses.

5 Illustrative Examples

5.1 Example 1: Simple Prerequisite Tree

Consider the prerequisites for MATH249, which requires satisfying two groups of requirements: $(\text{MATH135} \vee \text{MATH145}) \wedge (\text{MATH136} \vee \text{MATH146})$.

- **Scenario 1:** If the algorithm first selects MATH135, it satisfies the first group. For the second group, if MATH136 is chosen (perhaps due to a higher weight or because MATH135 is a prerequisite), the final path is $\{\text{MATH135}, \text{MATH136}\}$.
- **Scenario 2:** If MATH145 is chosen for the first group, then MATH146 must be chosen for the second, as MATH136 typically requires MATH135. The resulting path is $\{\text{MATH145}, \text{MATH146}\}$.

The algorithm selects the path with the optimal total cost, considering both course weights and dependencies.

5.2 Example 2: Complex Prerequisite Tree

The prerequisites for STAT330 involve multiple nested requirements:

- **STAT330:** requires $\text{MATH237} \wedge (\text{STAT230} \vee \text{STAT240}) \wedge \text{STAT231}$.
- **MATH237:** requires $(\text{one of } [\text{MATH106}, \text{MATH114}, \text{MATH115}, \text{MATH136}, \text{MATH146}]) \wedge (\text{one of } [\text{MATH128}(\geq 70), \text{MATH138}(\geq 60), \text{MATH148}])$.
- **STAT230:** requires $(\text{one of } [\text{MATH116}, \text{MATH118}, \text{MATH128}]) \wedge (\text{one of } [\text{MATH137}(\geq 80), \text{MATH138}])$.
- **STAT231:** requires one of $[\text{MATH118}(\geq 70), \text{STAT220}(\geq 70), \text{STAT230}]$.

A possible path that minimizes the number of unique courses could be constructed as follows:

1. To satisfy MATH237, select MATH136 and MATH138.
2. To satisfy STAT230, select MATH116 and reuse the already selected MATH138.
3. To satisfy STAT231, reuse the now-selected STAT230.

This yields the minimal path: $\{\text{MATH136}, \text{MATH138} (\geq 60), \text{MATH116}, \text{STAT230} (\geq 60), \text{MATH237},$
The algorithm systematically avoids redundant courses by considering the global set of selected courses.

5.3 Example 3: Node Weight Calculation

Let's calculate the weight for ACTSC431, given ratings (`liked` = 68.42, `easy` = 28.57, `useful` = 100), number of ratings = 19, and a user preference for "likeness."

- $\beta = (0.7, 0.15, 0.15)$.
- The reliability scalar λ is 0.9 (since $r_i = 19 < 50$).
- The weighted score is $s_{\text{ACTSC431}} = 0.7(68.42) + 0.15(28.57) + 0.15(100) = 67.18$.
- The final weight is $y_{\text{ACTSC431}} = 0.9 \times 67.18 = 60.46$.

If a course has all null ratings, assuming a global minimum rating of 18.18 and $\gamma = 0.5$:

- The imputed scores would be $x_1 = x_2 = x_3 = 9.09$.
- With $\lambda = 0.97$, the resulting weight y_i would be significantly lower, correctly flagging the course as a high-risk or unknown-quality option.